

Digitally Distributed Orchestra: DDO

Andrea Compton, Ashley Hedberg, and Tyler Lubeck

December 5, 2014

Contents

1	Project Goals	1
2	Final Design	1
2.1	Client-Server Relationship	1
2.2	Sheet Music Representation	2
2.3	Audio Generation	4
2.4	Design Diagram	4
3	Analysis and Reflection	4
3.1	Deliverable Status	4
3.2	Best Decision	4
3.3	Worst Decision	4
4	Abstractions and Language	4
5	Implementation Strategy	5
5.1	Timeline	5
5.2	Reflection on Division of Labor	5
6	Bug Report	5
7	Code Overview	5
8	How To Run	5

1 Project Goals

Our goal for this project was to create a digitally distributed orchestra. Given a representation of a song (i.e. digitized sheet music), one computer would assign various client machines a note to play. The client machines would then play the song together.

At a minimum, we wanted our orchestra to be able to play a hard-coded sequence of frequencies and durations in unison. Any client with [PyAudio](#) installed could serve as a client in our minimum deliverable.

One stretch goal was being able to play any song written in a specific format. This would involve handling songs which played notes in multiple octaves as well as those containing sharps and/or flats. Another stretch goal was supporting more types of client machines, such as Windows and Linux machines without the PyAudio package.

2 Final Design

2.1 Client-Server Relationship

For this project, we need to implement some notion of clients communicating with a central server. We considered two options for this.

2.1.1 Option 1: Twisted

The first was using the [Twisted](#) Python package. Given that parsing of the song file and playing the audio must be done in Python, using this option would allow the entire project to be implemented in one language. Twisted has a built-in notion of connections being made and lost, which would allow our server to more easily keep track of clients and distribute song notes accordingly. It also has a notion of executing tasks at specific time deltas, which would help enforce a consistent concept of time among the clients. However, none of us had used this framework before, and it involves some fairly low-level asynchronous calls, which had the potential to be challenging.

2.1.2 Option 2: Erlang

The second was implementing the server/client relationship in Erlang and the file/audio parsing in Python. This option would allow us to take advantage of Erlangs built-in message passing for note distribution. File parsing and sound generation would be more easily accomplished in Python, and we could use the [ErlPort](#) package to communicate between the two languages. However, we would have to implement the concept of client connections being lost ourselves.

2.1.3 Decision

We opted to use the first method. Being able to redistribute notes when a client drops seems like an important requirement for a digitally distributed orchestra. This functionality is built

in to the Twisted servers, whereas we would have to implement it ourselves in the Erlang clients.

2.2 Sheet Music Representation

2.2.1 Inspiration from Guitar Tabs

Our initial text-based representation of a song was the following:

```
Song Title
Beats per minute (an integer)
Note 1: ---X---X---
Note 2: --XX---X---
...
```

Note 1, 2, etc. will be A, B, and so on. A dash represents a beat where the given note does not play, and an X represents a beat where the given note does play.

2.2.2 Refinement of Long Note Representation

After we began implementing the parser for this format, we realized that this was not the best representation. We modified our song format to include the dash character (-), numbers, and the letter h. A dash represents one beat of silence. Numbers represent the number of times the note plays in one beat. This allows us to represent quarter notes, eighth notes, and so on. The h indicates a hold of the previous note, and allows us to represent notes lasting longer than one beat. "Hot Cross Buns," for example, looked like this in the new format:

```
hot cross buns
60
B:1--1----1---
A:-1----22-1--
G:--1h22----1h
```

We then discovered that handling held notes with this format was cumbersome. If a 1 was found while parsing, it was difficult to discern whether the note played for one beat or was supposed to be held over multiple beats. As a result, we again modified our format so that all beats over which a note was supposed to be held were designated by the h character. This led us to the following representation of "Hot Cross Buns":

```
hot cross buns
60
B:1--1----1---
A:-1----22-1--
G:--hh22----hh
```

2.2.3 Handling of Sharps and Flats

We chose to represent sharps and flats using the # and b characters, respectively.

2.2.4 Handling of Octaves

The octave for each note is designated by a number (0-8) placed after the letter name. Middle C is represented as C4, and all other notes are offset from that (i.e., the B below middle C is B3, whereas the B above middle C is B4).

If a conductor decided he/she wanted to play "Hot Cross Buns" using A# instead of A and up one octave from middle C, the song would be represented as follows:

```
hot cross buns
60
B5:1--1----1---
A#5:-1-----22-1--
G5:--hh22----hh
```

Note that whitespace is allowed in note names to allow for beat alignment in the song files.

2.3 Audio Generation

We chose to use the PyAudio package for this task. PyAudio provides Python bindings for [PortAudio](#), a cross-platform audio I/O library. Using this package allowed our orchestra to support any computer architecture without implementing architecture-specific logic ourselves.

We also considered using system calls on Linux machines and the [winsound](#) module on Windows machines as fallbacks if PyAudio was not installed. However, because the EECS staff installed PyAudio on all of the Halligan lab machines for us, this was not necessary to test our orchestra with a large number of machines.

2.4 Design Diagram

3 Analysis and Reflection

3.1 Deliverable Status

Our orchestra can play any song provided in the format described earlier in this report. Any machine with PortAudio and PyAudio can join as a musician. The conductor does not need to have these packages installed. We also accomplished an extra stretch goal of having the musician's terminal light up when it is playing a note.

3.2 Best Decision

3.3 Worst Decision

4 Abstractions and Language

We have four major abstractions:

- The "conductor," or the server with which the musicians communicate.
- The "musician," or a client which is given a note to play.
- The parser, which takes a text-based representation of a song and converts it into a set of frequencies, each of which has a list of starting times and durations, that the conductor can distribute to musicians.
- The audio module, which attempts several different methods of playing sound based on what packages are available on the client machine.

5 Implementation Strategy

5.1 Timeline

- November 10:
 - Have a working connection between conductor and musicians
 - Handle new musicians and dropped connections
 - Generate at least one test file in parsable format
- November 17:
 - Gain access to Halligan 105 for testing
 - Request that packages be installed on Linux machines and Mac lab machines
 - Be able to parse input file through the conductor and send sets of notes to musicians
- November 24:
 - Musicians can play their specified note
 - Enforce consistent sense of timing among machines
- December 1:
 - Have musicians attempt to reconnect to conductor when a connection is lost

- Light up terminal screen when musician is actively playing
- December 5:
 - Have conductors and musicians exit more gracefully when the song is over

5.2 Reflection on Division of Labor

We originally planned to work on all aspects of this project as a group. However, scheduling issues prevented this from happening in several cases. As a result, not all group members were present for certain design changes. More clearly defining who was responsible for which portions of the implementation may have resulted in a more refined finished product.

6 Bug Report

7 Code Overview

8 How To Run