Digitally Distributed Orchestra: DDO

Andrea Compton, Ashley Hedberg, and Tyler Lubeck

December 5, 2014

Contents

1 Project Goals

Our goal for this project was to create a digitally distributed orchestra. Given a representation of a song (i.e. digitized sheet music), one computer would assign various client machines

a note to play. The client machines would then play the song together.

At a minimum, we wanted our orchestra to be able to play a hard-coded sequence of frequencies and durations in unison. Any client with PyAudio installed could serve as a client in our minimum deliverable.

One stretch goal was being able to play any song written in a specific format. This would involve handling songs which played notes in multiple octaves as well as those containing sharps and/or flats. Another stretch goal was supporting more types of client machines, such as Windows and Linux machines without the PyAudio package.

2 Final Design

2.1 Client-Server Relationship

For this project, we need to implement some notion of clients communicating with a central server. We considered two options for this.

2.1.1 Option 1: Twisted

The first was using the Twisted Python package. Given that parsing of the song file and playing the audio must be done in Python, using this option would allow the entire project to be implemented in one language. Twisted has a built-in notion of connections being made and lost, which would allow our server to more easily keep track of clients and distribute song notes accordingly. It also has a notion of executing tasks at specific time deltas, which would help enforce a consistent concept of time among the clients. However, none of us had used this framework before, and it involves some fairly low-level asynchronous calls, which had the potential to be challenging.

2.1.2 Option 2: Erlang

The second was implementing the server/client relationship in Erlang and the file/audio parsing in Python. This option would allow us to take advantage of Erlangs built-in message passing for note distribution. File parsing and sound generation would be more easily accomplished in Python, and we could use the ErlPort package to communicate between the two languages. However, we would have to implement the concept of client connections being lost ourselves.

2.1.3 Decision

We opted to use the first method. Being able to redistribute notes when a client drops seems like an important requirement for a digitally distributed orchestra. This functionality is built

in to the Twisted servers, whereas we would have to implement it ourselves in the Erlang clients.

2.2 Sheet Music Representation

2.2.1 Inspiration from Guitar Tabs

Our initial text-based representation of a song was the following:

```
Song Title
Beats per minute (an integer)
Note 1: ---X---X---
Note 2: --XX---X---
...
```

Note 1, 2, etc. will be A, B, and so on. A dash represents a beat where the given note does not play, and an X represents a beat where the given note does play.

2.2.2 Refinement of Long Note Representation

After we began implementing the parser for this format, we realized that this was not the best representation. We modified our song format to include the dash character (−), numbers, and the letter h. A dash represents one beat of silence. Numbers represent the number of times the note plays in one beat. This allows us to represent quarter notes, eighth notes, and so on. The h indicates a hold of the previous note, and allows us to represent notes lasting longer than one beat. "Hot Cross Buns," for example, looked like this in the new format:

```
hot cross buns
60
B:1--1----1---
A:-1----22-1--
G:--1h22----1h
```

We then discovered that handling held notes with this format was cumbersome. If a 1 was found while parsing, it was difficult to discern whether the note played for one beat or was supposed to be held over multiple beats. As a result, we again modified our format so that all beats over which a note was supposed to be held were designated by the h character. This led us to the following representation of "Hot Cross Buns":

```
hot cross buns
60
B:1--1----1---
A:-1----22-1--
G:--hh22----hh
```

2.2.3 Handling of Sharps and Flats

We chose to represent sharps and flats using the # and b characters, respectively.

2.2.4 Handling of Octaves

The octave for each note is designated by a number (0-8) placed after the letter name. Middle C is represented as C4, and all other notes are offset from that (i.e., the B below middle C is B3, whereas the B above middle C is B4).

If a conductor decided he/she wanted to play "Hot Cross Buns" using A# instead of A and up one octave from middle C, the song would be represented as follows:

```
hot cross buns
60
 B5:1--1----1---
A#5:-1----22-1--
 G5:--hh22----hh
```

Note that whitespace is allowed in note names to allow for beat alignment in the song files.

2.3 Audio Generation

We chose to use the PyAudio package for this task. PyAudio provides Python bindings for PortAudio, a cross-platform audio I/O library. Using this package allowed our orchestra to support any computer architecture without implementing architecture-specific logic ourselves.

We also considered using system calls on Linux machines and the winsound module on Windows machines as fallbacks if PyAudio was not installed. However, because the EECS staff installed PyAudio on all of the Halligan lab machines for us, this was not necessary to test our orchestra with a large number of machines.

2.4   Design Diagram
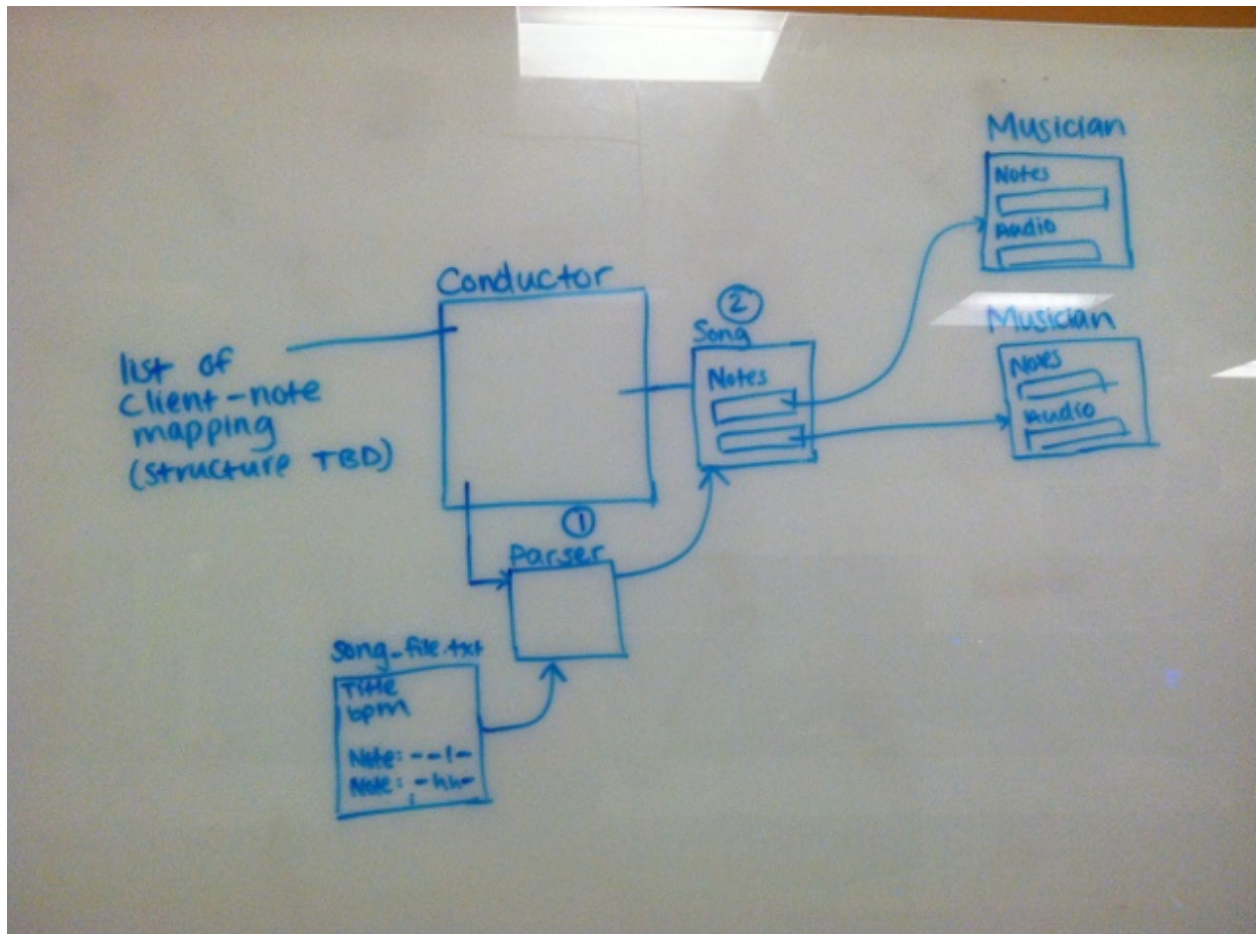
3   Analysis and Reflection

3.1   Deliverable Status

Our orchestra can play any song provided in the format described earlier in this report. Any machine with PortAudio and PyAudio can join as a musician. The conductor does not need to have these packages installed. We also accomplished an extra stretch goal of having the musician's terminal light up when it is playing a note.

3.2   Best Decision

There were many design decisions that shaped our project, but the best decision we made was to create a separate parsing module. Having the conductor call the module with the given song allows us to parse the song however we choose. This modularity gives us the flexibility

to change the song format in the future, and it allowed us to accomplish our stretch goals of playing sharps and flats as well as notes in multiple octaves without redesigning our entire project.

3.3 Worst Decision

The worst decision we made was to connect to the sound card every time a note was played. Connecting to and disconnecting from the sound card causes a horrid clicking noise, ruining the fluidity of the song. Changing this, however, would require a significant redesign of our musician. This may have been caused by our decision to use PyAudio to play a sine wave, which stopped abruptly after the specified duration ended. It does not resemble musical instruments, which gradually get softer as the note finishes.

4 Abstractions and Language

We have three major abstractions:

- The "conductor," or the server with which the musicians communicate.

- The "musician," or a client which is given a note to play.

- The parser, which takes a text-based representation of a song and converts it into a set of frequencies, each of which has a list of starting times and durations, that the conductor can distribute to musicians.

The implementation also contains an audio module. This module could be extended to attempt several different methods of playing sound based on what packages are available on the client machine. At present, the module either plays sound using PyAudio or does not play sound at all.

5 Implementation Strategy

5.1 Timeline

- November 10:
  - Had a working connection between conductor and musicians
  - Handled new musicians and dropped connections
  - Generated test files in parsable format
  - Parsed input file through the conductor and sent sets of notes to musicians
- November 17:
  - Gained access to Halligan 105 for testing
  - Requested that packages be installed on Linux machines and Mac lab machines
  - Began writing PyAudio-based play function

- November 24:
  - Abandoned attempts to prevent the sound card of clicking
  - Musicians could play their specified note
- December 1:
  - Enforced consistent sense of timing among machines
  - Had musicians attempt to reconnect to conductor when a connection is lost
  - Lit up terminal screen when musician is actively playing
- December 5:
  - Had musicians exit gracefully if they lose their connection to the conductor

5.2   Reflection on Division of Labor

We originally planned to work on all aspects of this project as a group. However, scheduling issues prevented this from happening in several cases. As a result, not all group members were present for certain design changes. More clearly defining who was responsible for which portions of the implementation may have resulted in a more refined finished product.

6   Bug Report

At one point, we encountered race conditions regarding accessing the sound card. If a musician was instructed to play two notes back-to-back, the first note would often still be playing when it came time for the second note to play. As a result, the second note could not access the sound card, and it would never be played.

We solved this problem by setting the duration of each note to 90% of its true value. The notes sounded correct to a human listener, and there was enough space in between sound card accesses that all notes played correctly.

7   Code Overview

Our code is available on GitHub. An explanation of the program structure follows.

7.1   Python Modules

audio.py Contains an audio class which constructs a function with which to play sound, based on packages available on the musician.

conductor.py Contains the implementation of the conductor.

musician.py Contains the implementation of the musician.

note_to_freq.py Contains logic to convert from a letter note (such as C4 or A#5) to a frequency. Also contains logic to calculate the duration of a note based on the BPM of the song.

parse.py Contains logic to concurrently parse a given song file.

suppress_errors.py Contains handlers necessary to avoid printing excessive output when PyAudio encounters errors.

7.2   Other Files

songs/* A collection of songs in the correct format.

note.txt An ASCII note that can be displayed when a musician plays.

requirements.txt A list of necessary packages that can be installed via pip.

7.3   How To Run

1. Install PortAudio for your machine's architecture.

2. Install PyAudio.

3. Clone the code from GitHub.

4. Install other Python packages:
   ```
   pip install -r ConcurrentMusic/requirements.txt
   ```

5. To start a conductor:
   ```
   python conductor.py --song <path to song> --port <port>
   ```
   The port number is optional and defaults to 8123.

6. To start a musician:
   ```
   python musician.py --ip <conductor IP> --port <conductor port>
   ```
   The port number is again optional and also defaults to 8123.